



Reactive Streams and Akka Streams

Markus Jura (@markusjura)
& Lutz Huehnken (@lutzhuehnken)



Streams

Common Use of Streams

- Bulk Data Transfer
- Real Time Data Sources
- Batch Processing of large data sets
- Monitoring and Analytics

What is a Stream?

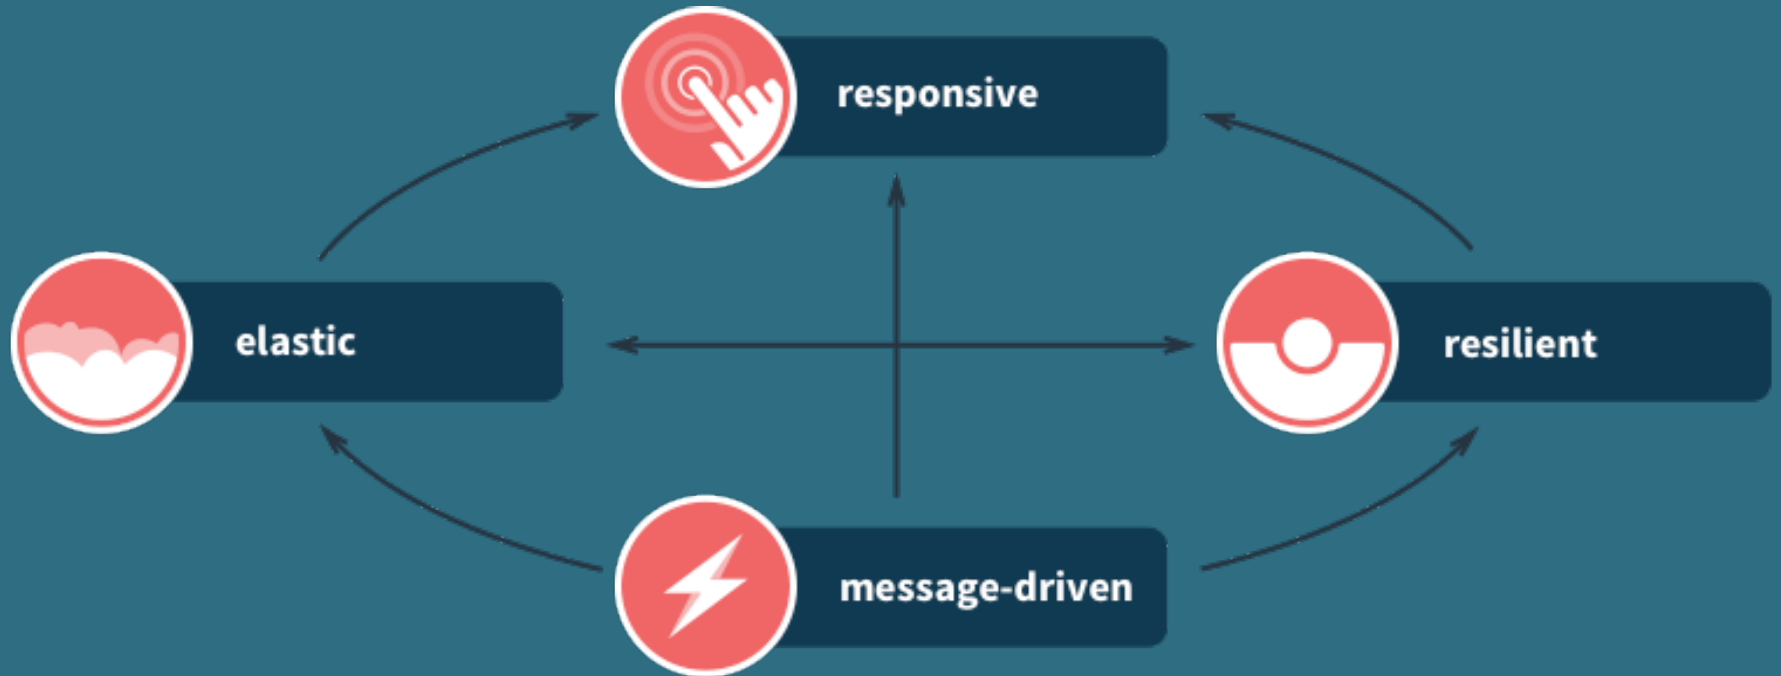
- Ephemeral flow of data
- Potentially unbounded in size
- Processed by describing transformation of data

Reactive Streams

Reactive Streams Projects / Companies

- Typesafe
 - Akka Streams
- Netflix
 - rxJava / rxScala
- Pivotal
 - Spring Reactor
- Redhat
 - Vert.x
- Oracle

Reactive




Reactive Streams Overview

- How do we:
 - Handle potentially infinite streams of data?
 - Handle data in a reactive manner?
 - Achieve asynchronous non-blocking data flow?
 - Avoid out of memory errors?


Supply and Demand

- Data Items Flow Downstream
- Demand Flows Upstream
- Data Items flow only when there is demand.



Dynamic Push-Pull

- “Push” behavior when consumer is faster
- “Pull” behavior when producer is faster
- Switches automatically between these
- Batching demand allows batching data



Reactive Streams Specification

- Interface Specification
 - Java interfaces for implementations
- TCK
 - Test Harness to validate implementations
- Specification Website
 - <http://www.reactive-streams.org/>

Publisher

```
package org.reactivestreams;  
public interface Subscriber<T> {...}  
public interface Publisher<T> {  
    public void subscribe<T>(  
        Subscriber<T> subscriber);  
}
```

Subscriber

```
public interface Subscriber<T> {  
    public void onSubscribe(  
        Subscription subscription);  
    public void onNext(T element);  
    public void onComplete();  
    public void onError(Throwable cause);  
}
```

Subscription

```
public interface Subscription {  
    public void cancel();  
    public void request(long elements);  
}
```

Backpressure

- Downstream consumers pushing back on the producer to prevent flooding.
- In reactive-streams:
 - Consumers stop requesting more elements until they are ready to process more.
 - Producers only fire elements if there is demand.

Why Backpressure?

- Explicitly design for system overload
 - Demand is propagated throughout the WHOLE flow
 - Can decide WHERE to handle overload
- Limit the number of in-flight messages throughout the system
 - Bounded memory consumption
 - Bounded cpu contention
- Recipient is in control of incoming data rate
- Data in flight is bounded by signaled demand

Buffering

- We can prefetch stream elements by requesting more than we really need.
- We can use this technique to ensure no "sputter" in the stream.
- We can also use this technique to pull faster than downstream consumer.

Akka Streams & Actors

Basic Actor

```
case class Greeting(who: String)

class GreetingActor extends Actor with
  ActorLogging {
  def receive = {
    case Greeting(who) => log.info("Hello " + who)
  }
}

val system = ActorSystem("MySystem")
val greeter: ActorRef =
  system.actorOf(Props[GreetingActor])

greeter ! Greeting("London Scala User Group")
```

Properties of Actors

- Message Based / Event Driven
- Isolated State
- Sane Concurrency Model
- Isolated Failure Handling (Supervision)

Akka Streams – A Bridge

- Converts Publisher/Subscriber API into Actor messages
- Simplify creating Publisher/Subscribers using Actors
- Attach Reactive streams into existing Akka applications

Creating an ActorSubscriber

```
import akka.stream._

class PrintlnActor extends Actor with ActorSubscriber {
  val requestStrategy = OneByOneRequestStrategy
  def receive: Receive = {
    case ActorSubscriberMessage.OnNext(element) =>
      println(element)
  }
}

val printlnActor:ActorRef =
  system.actorOf(Props[PrintlnActor], "println")

val subscriber = ActorSubscriber(printlnActor)
```

Creating an ActorPublisher

```
import akka.stream._

class IntActor extends Actor with ActorPublisher[Int] {
  def receive: Receive = {
    case ActorPublisherMessage.Request(elements) =>
      while (totalDemand > 0) { onNext(1) }
  }
}

val intActor: ActorRef =
  system.actorOf(Props[IntActor], "intActor")

val publisher = ActorPublisher(intActor)
```


Why use Actors as Publishers?

- Actors are smart
 - They can keep internal state to track demand and supply
 - They can buffer data to meet anticipated demand
- Actors are powerful
 - They can spin up child actors to meet demand
 - With Akka clustering, can spread load across multiple machines
- Actors are resilient
 - On exception, actor can be killed and restarted by supervisor
 - Actor interaction is thread-safe, and actor state is private


Actors Demo

Flows

Linear Transformations



Linear Transformations




Linear Stream Transformations

- Deterministic (like for collections)
 - map, filter, collect, grouped, drop, take, groupBy, ...
- Time-Based
 - takeWithin, dropWithin, groupedWithin, ...
- Rate-Detached
 - expand, conflate, buffer, ...
- asynchronous
 - mapAsync, mapAsyncUnordered, ...

Non-linear Stream Transformations


- Fan-In
 - merge, concat, zip
- Fan-Out
 - broadcast, route, unzip

Fan In



merging the data means splitting the demand

Fan Out



splitting the data means merging the demand

Materialization

- Akka Streams separate the *what* from the *how*
 - declarative Source/Flow/Sink DSL to create blueprint
 - FlowMaterializer turns this into running Actors
- this allows alternative materialization strategies
 - optimization
 - verification / validation
 - cluster deployment
- only Akka Actors for now, but more to come!

FlowGraph Demo

Play

Streaming Data

- Play can stream data using Iteratees and Enumerators
- Streamed data through chunked encoding, using `Ok.stream()`
- But Iteratees and Enumerators are complicated.
- And Reactive Streams are simple.

Play 2.4 Experimental Features

- Reactive Streams Integration!
- Adapts Futures, Promises, Enumerators and Iteratees
- Note: potential for event loss, no performance tuning
- All access through `play.api.libs.streams.Streams`

Streaming video through Play

```
def stream = Action {  
  val headers =  
    Seq(CONTENT_TYPE -> "video/mp4",  
        CACHE_CONTROL -> "no-cache")  
  val framePublisher =  
    video.FFMpeg.readFile(mp4, Akka.system)  
  val frameEnumerator =  
    Streams.publisherToEnumerator(framePublisher)  
  val bytesEnumeratee = Enumeratee.map[Frame](encodeFrame)  
  val chunkedVideoStream =  
    frameEnumerator.through(bytesEnumeratee)  
  Ok.stream(chunkedVideoStream).withHeaders(headers: _*)  
}
```

Play Demo

Questions

Thanks

@markusjura

@lutzhuehnken





©Typesafe 2014 – All Rights Reserved